MySQL Performance

Sang Shin www.javapassion.com "Learning is fun!"



Topics

- Query cache
 - > Overview
 - > System variables
- Indexing
 - > Overview
 - Example query
 - > Strategy
 - > Runtime counters
- EXPLAIN
- Slow query log

Query Cache: Overview

What is Query Cache?

- The query cache stores the text of a SELECT statement together with the corresponding result that was sent to the client
 - If an identical statement (case sensitive) is received later, the server retrieves the results from the query cache rather than parsing and executing the statement again
- The query cache is shared among sessions, so a result set generated by one client can be sent in response to the same query issued by another client

Why Query Cache?

- The query cache is extremely useful in an environment where you have tables that do not change very often and for which the server receives many identical queries
 - This is a typical situation for many Web servers that generate many dynamic pages based on database content

How Does Query Cache Work?

- When a table is changed, any relevant entries in the query cache are discarded
 - > Prevents the query cache from being stale
- The query has to be exactly the same including the case
 - "SELECT * from person" is considered a different query from "select * from person"

Query Cache: System Variables

Configuration variables - You can set

- query_cache_size
 - The amount of memory allocated for caching query results. The default value is 0, which disables the query cache.
- query_cache_type
 - > 0 or OFF: Query caching is turned off
 - > 1 or ON: Query caching is turned on
 - > 2 or DEMAND: Query caching is done on demand basis
- query_cache_limit
 - Maximum result set size. Result sets larger than this value will not be cached

Configuration variables

```
mysql> SHOW variables LIKE '%query cache%';
Variable name
                             | Value
have query cache
query cache limit
                               1048576
query cache min res unit
                               4096
query cache size
                               1048576
query cache type
                               ON
query cache wlock_invalidate | OFF
6 rows in set (0.00 sec)
```

Runtime counter variables

- Qcache inserts
 - > Total number of times queries are cached
- Qcache queries in cache
 - Number of queries that are currently present in the query cache
 - > Will be reset to 0 when query cache is flushed
- Qcache_hits
 - > Number of times queries in the cache are hit
 - > High number of this indicates that the cache is being used effectively

Runtime counter variables

```
mysql> SHOW STATUS LIKE 'Qcache%';
Qcache free_blocks
Qcache free memory
Qcache hits
Qcache inserts
Qcache lowmem prunes
Qcache not cached
Qcache queries in cache
Qcache_total_blocks | 0 |
-----+
8 rows in set (0.00 sec)
```

Counter variables after 1st query

```
mysql> select * from employees where employee id = 5;
+----+
employee_id | ename | d_id | salary
        5 | jones | 3 | 5000.00 |
mysql> SHOW STATUS LIKE 'Qcache%';
Qcache free blocks
Qcache_free_memory | 1038376
Qcache hits
Qcache inserts
Qcache_lowmem_prunes | 0
Qcache not cached
Qcache queries in cache | 1
Qcache total blocks
```

Counter variables after 2nd Identical query

```
mysql> select * from employees where employee id = 5;
employee id | ename | d id | salary
           5 | jones | 3 | 5000.00 |
mysql> SHOW STATUS LIKE 'Qcache%';
Variable_name | Value
Qcache free blocks
Qcache free memory | 1038376 |
Qcache hits
                                    <--- Cache is hit
Qcache inserts
Qcache lowmem prunes | 0
Qcache not cached
Qcache queries in_cache | 1
Qcache total blocks
```

Selective Caching

```
// Using SQL NO CACHE means "don't cache this guery" when
// query cache type is set to DEMAND
mysql> select SQL NO CACHE * from employees where employee id = 1;
+----+
| employee_id | ename | d_id | salary |
+----+
 1 | jack | 1 | 3000.00 |
+----+
1 row in set (0.00 sec)
// Using SQL CACHE means "cache this query" when
// query cache type is set to DEMAND
mysql> select SQL CACHE * from employees where employee id = 1;
+----+
| employee_id | ename | d id | salary |
+----+
 1 | jack | 1 | 3000.00 |
+----+
1 row in set (0.00 sec)
```

Indexing: Overview

What is and Why Indexing?

- A logically separate table that contains sorted values of a column
 - You can have multiple indexes for multiple columns of a table
- Prevents a full table scan
 - Without indexing, a full table scan might be needed to find a match
- Inserting or deleting a row from an indexed column is more efficient
 - Since you do not have to move around the rows with full number of columns

Indexing in Joined Query

- Indexes are even more valuable when you're running queries involving joins on multiple tables
 - In a single-table query, the number of values you need to examine per column is the number of rows in the table
 - In a multiple-table query, the number of possible combinations exponentially increases because it's the product of the number of rows in all tables involved

Where Indexing is Used in MySQL

- To speed up searches for rows matching terms of a WHERE clause or rows that match rows in other tables when performing joins
- For queries that use the MIN() or MAX()
 functions, the smallest or largest value in an
 indexed column can be found quickly
 without examining every row.
- MySQL can often use indexes to perform sorting and grouping operations quickly for ORDER BY and GROUP BY clauses

Indexing: Example Query

Example Query Case

- Suppose we have 3 tables, t1, t2, t3
 - > Each table contains a column c1, c2, and c3
 - > Each table contains 100 rows
- Suppose we have the following query
 - > SELECT t1.i1, t2.i2, t3.i3
 - > FROM t1, t2, t3
 - > WHERE t1.c1 = t2.c2 AND t2.c1 = t3.c3;

When tables are NOT Indexed

- The possible number of rows to scan is
 - > 100 * 100 * 100 = 1 million
- It is because there is no knowing which row of a table contains which value without scanning them all. So in order to find the one that matches the WHERE clause, you might have to read all rows of all three tables
 - > 100 * 100 * 100 = 1 million

When tables are Indexed

- The number of rows to scan is
 - > 100 for the first table, t1, and 1 in the 2nd table, t2, and 1 in the 3rd table, t3
- Using indexes of t2 and t3, for each row of t1, you can find quickly find which one in t2 and t3 are matches

Indexing: Strategy

Which columns to Index?

- Index columns that you use for searching, sorting, or grouping
 - > Columns that appear in WHERE clause
 - Columns named in JOIN clauses
 - Columns that appear in ORDER BY or GROUP BY clauses
- Do not index columns that appear only in the output column list following the SELECT keyword

Example: Indexing Strategy

```
SELECT
c1 /* t1.c1 is not a good candidate for indexing */
FROM
t1 LEFT JOIN t2
ON t1.c2 = t2.c3 /* t1.c2 and t2.c3 are good candidate for indexing */
WHERE
c4 = 2000; /* t1.c4 is a good candidate for indexing */
```

Column Cardinality

- The cardinality of a column is the number of distinct values that it contains
 - > For example, a column that contains the values 1, 3, 7, 4, 7, 3 has a cardinality of four (1, 3, 4, 7).
- Indexes work best for columns that have a high cardinality relative to the number of rows in the table (that is, columns that have many unique values and few duplicates).

Indexing: Runtime counters

Indexing Runtime Counters

- Handler read first:
 - Number of times the first entry was read from an index. If this is high, it suggests that the server is doing a lot of full index scans.
- Handler_read_key:
 - Number of requests to read a row based on a key. If this is high, it is a good indication that your queries and tables are properly indexed.
- Handler read rnd:
 - Number of requests to read a row based on a fixed position. This will be high if you are doing a lot of queries that require sorting of the result.

EXPLAIN

What is EXPLAIN?

- Used to obtain information about how MySQL executes a SELECT statement
- When you precede a SELECT statement with the keyword EXPLAIN, MySQL displays information from the optimizer about the query execution plan. That is, MySQL explains how it would process the SELECT, including information about how tables are joined and in which order.

What is EXPLAIN?

- With the help of EXPLAIN, you can see where you should add indexes to tables to get a faster SELECT that uses indexes to find rows.
- You can also use EXPLAIN to check whether the optimizer joins the tables in an optimal order.

EXPLAIN Fields

- table
 - Indicates which table the output is about (for when you join many tables in the query)
- type
 - Indicates which type of join is being used. From best to worst the types are: system, const, eq_ref, ref, range, index, all
- possible_keys
 - Indicates possible indexes that can be used by MySQL
- key
 - Indicates which index (among the ones in possible keys) is actually used

EXPLAIN Fields (Continued)

- key_len
 - > The length of the key used. The shorter the better.
- ref
 - > Tells which column, or a constant, is used
- rows
 - Number of rows MySQL believes it must examine to get the data
 - > With indexing, smaller number is expected
- extra Extra info
 - The bad ones to see here are "using temporary" and "using filesort"

EXPLAIN for non-indexed Query

```
id: 1
select type: SIMPLE
   table: employees
    type: ALL
                                     Note that without indexing, the number
possible keys: d id
                                     of rows MySQL is expected to read is
    key: NULL
                                     the total number of rows of each table,
  key len: NULL
    ref: NULL
                                     5 for the employees table and 4 for the
    rows: 5
                                     departments table.
   Extra: Using where
id: 1
select type: SIMPLE
   table: departments
    type: ALL
possible keys: PRIMARY
    key: NULL
  key len: NULL
    ref: NULL
    rows: 4
   Extra: Using where; Using join buffer
2 rows in set (0.00 sec)
```

EXPLAIN for Indexed Query

```
id: 1
select type: SIMPLE
   table: employees
    type: ref
                                    Note that with indexing, the number of
possible keys: ename,d id
                                    rows MySQL is expected to read is 1 for
    key: ename
  key len: 257
                                    both employees and departments
    ref: const
                                    tables.
    rows: 1
   Extra: Using where
id: 1
select type: SIMPLE
   table: departments
    type: eq ref
possible keys: PRIMARY
    key: PRIMARY
  key_len: 4
    ref: mydb.employees.d id
    rows: 1
   Extra:
2 rows in set (0.05 sec)
```

Slow Query Log

Slow Query Log

- MySQL has a feature that allows you to log slow running queries to a file
- These queries are the candidates for optimization

Configuring Slow Query Log

```
mysql> SHOW VARIABLES LIKE '%long query%';
  -----+
| Variable_name | Value
| long_query_time | 10.000000 |
+----+
1 row in set (0.00 sec)
mysql> SET GLOBAL slow query log=ON;
Query OK, 0 rows affected (0.13 sec)
mysql> SHOW VARIABLES LIKE '%slow%';
| Variable_name | Value
log_slow_queries | ON
slow launch time | 2
slow query log | ON
slow query log file | C:\ProgramData\MySQL\MySQL Server 5.1\Data\ShinLaptop-slow.log |
4 rows in set (0.03 sec)
```

Configuring Slow Query Log

```
C:\Users\sang>type "C:\ProgramData\MySQL\MySQL Server 5.1\Data\ShinLaptop-slow.log"
C:\Program Files (x86)\MySQL\MySQL Server 5.1\bin\mysqld, Version: 5.1.44-community
(MySQL Community
Server (GPL)). started with:
TCP Port: 3306, Named Pipe: MySQL
Time
              Id Command Argument
# Time: 100329 15:15:08
# User@Host: root[root] @ localhost [127.0.0.1]
# Query time: 0.040000 Lock_time: 0.000000 Rows_sent: 0 Rows_examined: 0
use mydb;
SET timestamp=1269890108;
DROP TABLE IF EXISTS employees;
# Time: 100329 15:15:09
# User@Host: root[root] @ localhost [127.0.0.1]
# Query time: 0.030000 Lock time: 0.000000 Rows sent: 0 Rows examined: 0
SET timestamp=1269890109;
DROP TABLE IF EXISTS departments;
# User@Host: root[root] @ localhost [127.0.0.1]
# Query time: 0.100000 Lock time: 0.000000 Rows sent: 0 Rows examined: 0
SET timestamp=1269890109;
```

Thank you!

Sang Shin
http://www.javapassion.com
"Learning is fun!"

